

빅바이오 part2-1,  
딥러닝실습 (tensorflow),

# Convolutional Neural Networks


2016-4-5  
한성국

- Little bit more about Convolution Neural Networks
- Convolution Neural Networks Tutorial of Tensorflow

# Goals of this Tutorial


- Network architecture, training and evaluation
- Start from this template.
- Highlights
  - Simpler Convolutional Neural Networks (CNNs).
  - Convolution, relu, maxpooling, local response normalization.
  - Moving average → Enhancement in prediction performance.
  - Implementation of a learning rate schedule.
  - Prefetching queue for input data (large-scale training).
  - Multi-GPU Training ( c.f. distributed tensorflow recently released).

# Convolutional Neural Networks




- 5 convolution layers + 2 pooling layers + 2 FC layers.
- Dropout.
- Rectified linear unit.
- Local response Normalization.
- [ILSVRC 2012 result](#)

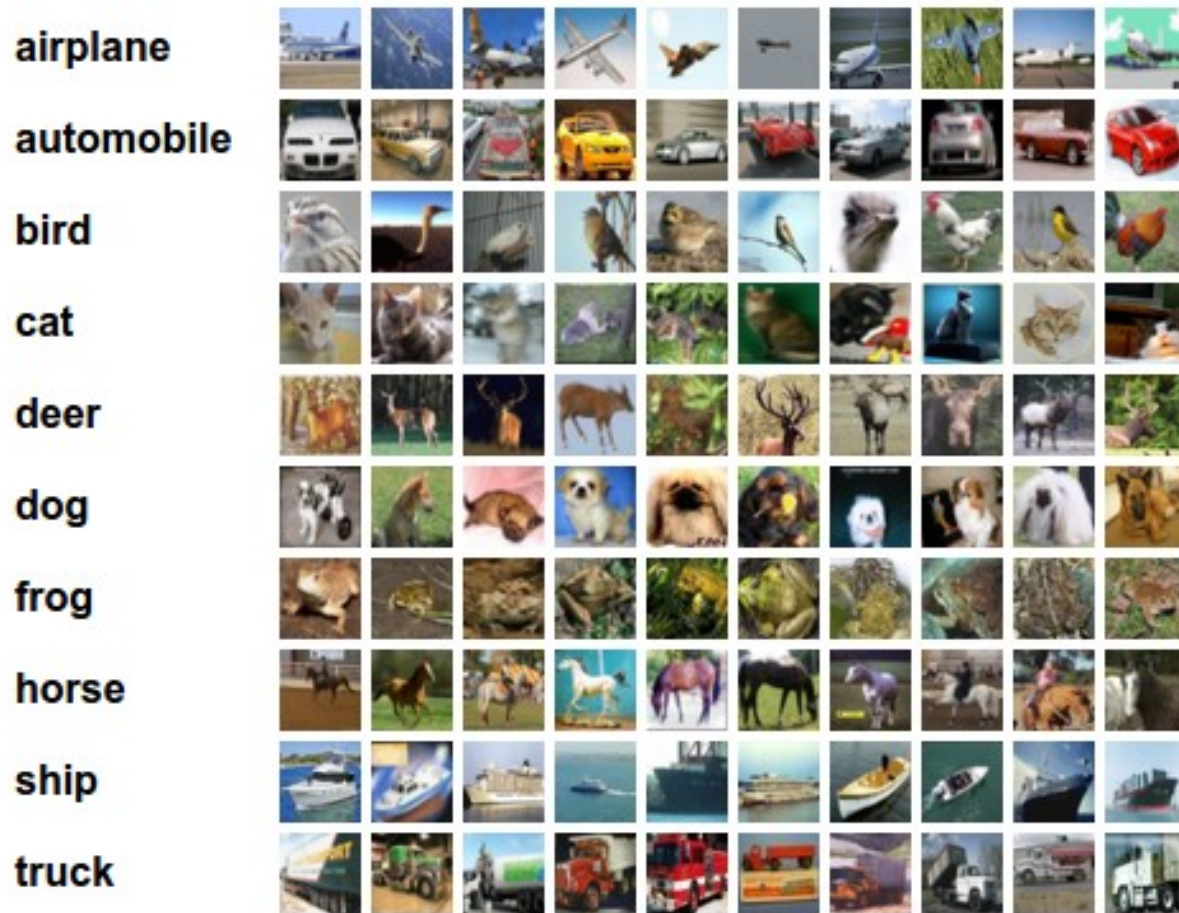
# Simpler CNNs



- # of Parameters  $\sim 1\text{M} \ll$  # of parameters Alexnet, googlenet, VGGnet.
- Operations: 19.5M.
- 86% prediction accuracy.



# CIFAR-10 data set



- 32x32pixels
- 10 classes
- Total dataset 60,000(Training :50000, Testing:10000)
- 6000/class

# Code template

| File                                    |   |
|---|---|
| <code>cifar10_input.py</code>           | Reads the native CIFAR-10 binary file format.             |
| <code>cifar10.py</code>                 | Builds the CIFAR-10 model.                                |
| <code>cifar10_train.py</code>           | Trains a CIFAR-10 model on a CPU or GPU.                  |
| <code>cifar10_multi_gpu_train.py</code> | Trains a CIFAR-10 model on multiple GPUs.                 |
| <code>cifar10_eval.py</code>            | Evaluates the predictive performance of a CIFAR-10 model. |

# Model inputs: Image preprocessing (cifar10\_input.py)


- Cropping
- Approx. Whitened to zero mean and unity std.
- Randomly Flip
- Image brightness
- Image contrast.



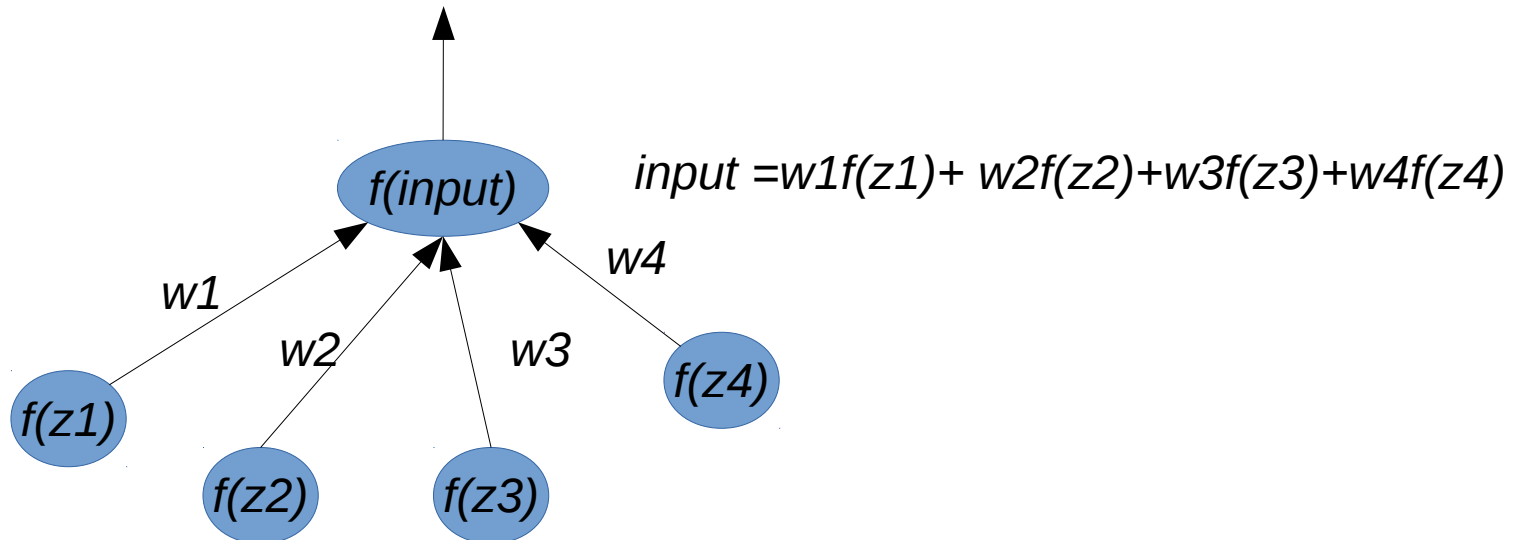
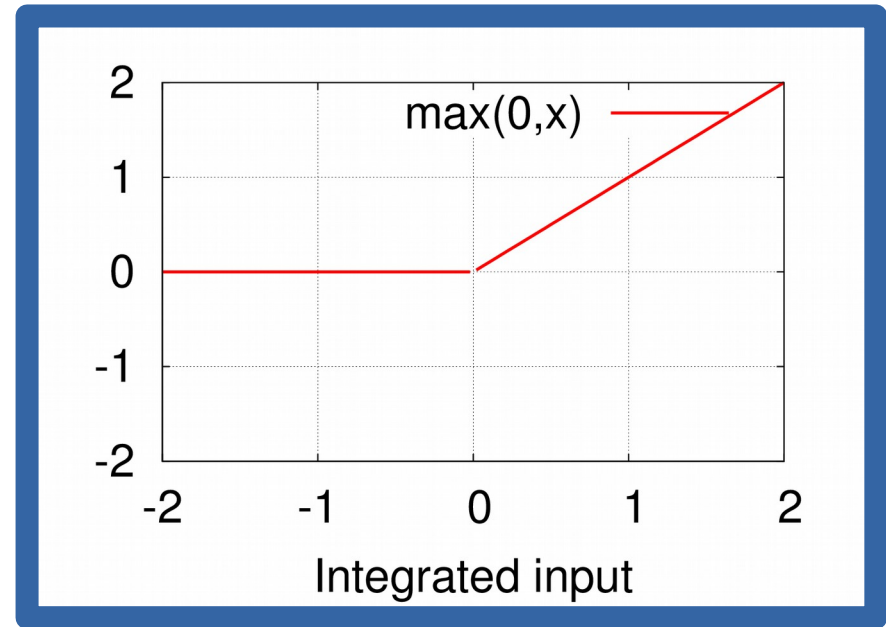
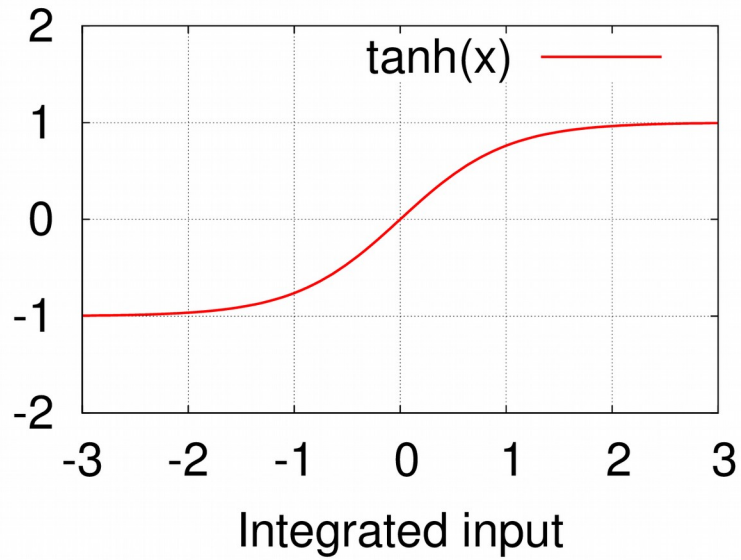
# Model Prediction (cifar10.py)

| Layer Name     |   |
|----------------|---|
| conv1          | Convolution, RELU.                          |
| pool1          | Max pooling(cf. Average pooling).           |
| norm1          | Local response normalization.               |
| conv2          | Convolution and Retified linear activation. |
| norm2          | Local response normalization.               |
| pool2          | Max pooling(cf. Average pooling).           |
| loca3          | Fully connected, RELU.                      |
| local4         | Fully connected, RELU.                      |
| softmax_linear | Exercise1: check to normalized predictions. |


Exercise2: change the FC structure to locally connected at the local3 and local4. Dropout.



# Neurons



# Local response normalization



$$b_{x,y}^i = a_{x,y}^i / \left( k + \alpha \sum_{j=\max(0, i-n/2)}^{\min(N-1, i+n/2)} (a_{x,y}^j)^2 \right)^\beta$$

- Neuron activations are normalized within the feature maps.
- Stable to quite large learning rate → Faster learning.

# EX1 and EX2

- Learning curve.
- Prediction accuracy.

# Model Training (cifar10\_train.py)

- Multinomial logistic regression (softmax)
- Loss function: cross-entropy
- Weight decay: L2 regularization
- Learning rate schedule: exponentially decays
- Loss Optimizer:
  - `tf.train.GradientDescentOptimizer`
  - `tf.train.AdagradOptimizer`
  - `tf.train.AdamOptimizer`

# Launching and Training the Model

```
$ python cifar10_train.py --help
```

optional arguments:

```
-h, --help                show this help message and exit
--batch_size BATCH_SIZE  Number of images to process in a batch.
--data_dir DATA_DIR     Path to the CIFAR-10 data directory.
--train_dir TRAIN_DIR    Directory where to write event logs and
```

checkpoint.

```
--max_steps MAX_STEPS    Number of batches to run.
--log_device_placement [LOG_DEVICE_PLACEMENT]
                          Whether to log device placement.
--nolog_device_placement
```

# Launching and Training the Model

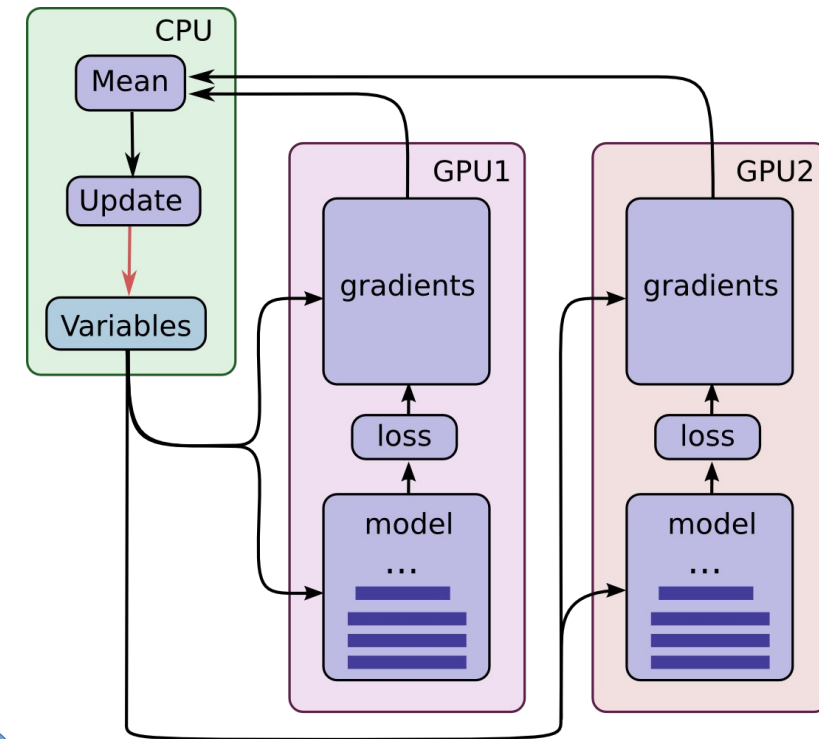

```
$ python cifar10_train.py
```

```
2016-04-02 00:03:50.972550: step 182840, loss = 0.27 (804.1 examples/sec; 0.159 sec/batch)
2016-04-02 00:03:52.596317: step 182850, loss = 0.22 (851.1 examples/sec; 0.150 sec/batch)
2016-04-02 00:03:54.144716: step 182860, loss = 0.30 (879.3 examples/sec; 0.146 sec/batch)
2016-04-02 00:03:55.798343: step 182870, loss = 0.29 (754.4 examples/sec; 0.170 sec/batch)
2016-04-02 00:03:57.381721: step 182880, loss = 0.26 (772.9 examples/sec; 0.166 sec/batch)
2016-04-02 00:03:58.935672: step 182890, loss = 0.35 (893.9 examples/sec; 0.143 sec/batch)
2016-04-02 00:04:00.487395: step 182900, loss = 0.29 (830.9 examples/sec; 0.154 sec/batch)
2016-04-02 00:04:02.218012: step 182910, loss= 0.29 (922.3 examples/sec; 0.139 sec/batch)
2016-04-02 00:04:03.771728: step 182920, loss = 0.34 (799.0 examples/sec; 0.160 sec/batch)
2016-04-02 00:04:05.343366: step 182930, loss = 0.27 (797.8 examples/sec; 0.160 sec/batch)
```

( GTX980Ti )

# Multiple GPU cards Implementation

- Model replica





\$python cifar10\_multigpu\_train.py --num\_gpus n

1 gtx980ti+1gtx960

```
sghan@aturing: ~/tensorflow-master/tensorflow/models/image/cifar10 116x
[2016-09-26 14:58:00] GPU 0 memory begin
o 0xcde320e67
: step 0, loss = 4.68 (13.2 examples/sec; 9.681 sec/bat

untime/gpu/pool_allocator.cc:244] PoolAllocator: After
ount=612093 evicted_count=1000 eviction_rate=0.00163374
n rate=0.00182484
: step 10, loss = 4.66 (909.0 examples/sec; 0.141 sec/b

step 20, loss = 4.64 (1023.9 examples/sec; 0.125 sec/batch)
: step 30, loss = 4.63 (928.2 examples/sec; 0.138 sec/batch)
: step 40, loss = 4.61 (845.4 examples/sec; 0.151 sec/batch)
: step 50, loss = 4.59 (949.6 examples/sec; 0.135 sec/batch)
: step 60, loss = 4.57 (953.3 examples/sec; 0.134 sec/batch)
: step 70, loss = 4.55 (828.5 examples/sec; 0.154 sec/batch)
: step 80, loss = 4.53 (818.1 examples/sec; 0.156 sec/batch)
: step 90, loss = 4.51 (878.3 examples/sec; 0.146 sec/batch)
: step 100, loss = 4.50 (801.3 examples/sec; 0.160 sec/batch)
: step 110, loss = 4.48 (816.1 examples/sec; 0.157 sec/batch)
: step 120, loss = 4.46 (848.6 examples/sec; 0.151 sec/batch)
: step 130, loss = 4.44 (839.9 examples/sec; 0.152 sec/batch)
: step 140, loss = 4.43 (877.6 examples/sec; 0.146 sec/batch)
: step 150, loss = 4.41 (940.3 examples/sec; 0.136 sec/batch)
: step 160, loss = 4.39 (753.7 examples/sec; 0.170 sec/batch)
: step 170, loss = 4.38 (821.0 examples/sec; 0.156 sec/batch)
: step 180, loss = 4.36 (914.7 examples/sec; 0.140 sec/batch)
: step 190, loss = 4.34 (877.5 examples/sec; 0.146 sec/batch)
: step 200, loss = 4.33 (896.4 examples/sec; 0.143 sec/batch)
: step 210, loss = 4.31 (876.2 examples/sec; 0.146 sec/batch)
: step 220, loss = 4.29 (920.7 examples/sec; 0.139 sec/batch)
: step 230, loss = 4.28 (883.5 examples/sec; 0.145 sec/batch)
: step 240, loss = 4.26 (946.6 examples/sec; 0.135 sec/batch)
: step 250, loss = 4.25 (862.0 examples/sec; 0.148 sec/batch)
: step 260, loss = 4.23 (948.4 examples/sec; 0.135 sec/batch)
: step 270, loss = 4.22 (881.7 examples/sec; 0.145 sec/batch)
```

1 gtx980ti

```
sghan@aturing: ~/tensorflow-master/tensorflow/models/image/cifar10 102x52
step 0, loss = 4.68 (11.0 examples/sec; 11.595 sec/batch)
step 10, loss = 4.66 (925.5 examples/sec; 0.138 sec/batch)
step 20, loss = 4.64 (872.3 examples/sec; 0.147 sec/batch)
step 30, loss = 4.61 (789.1 examples/sec; 0.162 sec/batch)
step 40, loss = 4.60 (751.7 examples/sec; 0.170 sec/batch)
step 50, loss = 4.59 (792.2 examples/sec; 0.162 sec/batch)
step 60, loss = 4.57 (957.6 examples/sec; 0.134 sec/batch)
step 70, loss = 4.55 (1014.2 examples/sec; 0.126 sec/batch)
step 80, loss = 4.53 (1022.2 examples/sec; 0.125 sec/batch)
step 90, loss = 4.51 (1027.0 examples/sec; 0.125 sec/batch)
step 100, loss = 4.48 (1013.8 examples/sec; 0.126 sec/batch)
step 110, loss = 4.48 (777.9 examples/sec; 0.165 sec/batch)
step 120, loss = 4.46 (857.1 examples/sec; 0.149 sec/batch)
step 130, loss = 4.44 (824.8 examples/sec; 0.155 sec/batch)
step 140, loss = 4.42 (798.4 examples/sec; 0.160 sec/batch)
step 150, loss = 4.41 (988.2 examples/sec; 0.130 sec/batch)
step 160, loss = 4.39 (744.8 examples/sec; 0.172 sec/batch)
step 170, loss = 4.37 (751.4 examples/sec; 0.170 sec/batch)
step 180, loss = 4.36 (763.3 examples/sec; 0.168 sec/batch)
step 190, loss = 4.34 (892.6 examples/sec; 0.143 sec/batch)
step 200, loss = 4.32 (712.3 examples/sec; 0.180 sec/batch)
step 210, loss = 4.31 (710.7 examples/sec; 0.180 sec/batch)
step 220, loss = 4.30 (824.7 examples/sec; 0.155 sec/batch)
step 230, loss = 4.27 (691.6 examples/sec; 0.185 sec/batch)
step 240, loss = 4.27 (723.0 examples/sec; 0.177 sec/batch)
step 250, loss = 4.25 (728.8 examples/sec; 0.176 sec/batch)
step 260, loss = 4.23 (743.4 examples/sec; 0.172 sec/batch)
step 270, loss = 4.22 (772.6 examples/sec; 0.166 sec/batch)
step 280, loss = 4.20 (677.5 examples/sec; 0.189 sec/batch)
step 290, loss = 4.18 (694.3 examples/sec; 0.184 sec/batch)
```

# References

- [Tensorflow.org](https://www.tensorflow.org)
- Vincent Dumoulin and Francesco Visin, “A guide to convolution arithmetic for deep learning”, arXiv:1603.07285 [stat.ML]
- Matthew D. Zeiler and Rob Fergus, “Visualizing and Understanding Convolutions Networks”, arXiv:1311.2901 [cs.CV]